

Hive Case Study :- E-Commerce Sales

Contents

1. Problem Statement

2. Steps Involved for Analysis :

- 2.1 Create the Key pair in AWS
- 2.2 Create S3 bucket and copy the csv data files into s3 bucket.
- 2.3 Create the EMR cluster and launch EMR that utilizes the Hive services
- 2.4 Move the data from the S3 bucket into the HDFS
- 2.5 Creating the database and launching Hive queries on EMR cluster
- 2.6 Creating the structure of database
- 2.7 Use partitioning and bucketing to run your queries as efficiently as possible
- 2.8 Run Hive queries to verify the analysis on given dataset.

3. Analysis Questions and Answers

4. Drop the database

5. Terminate the cluster

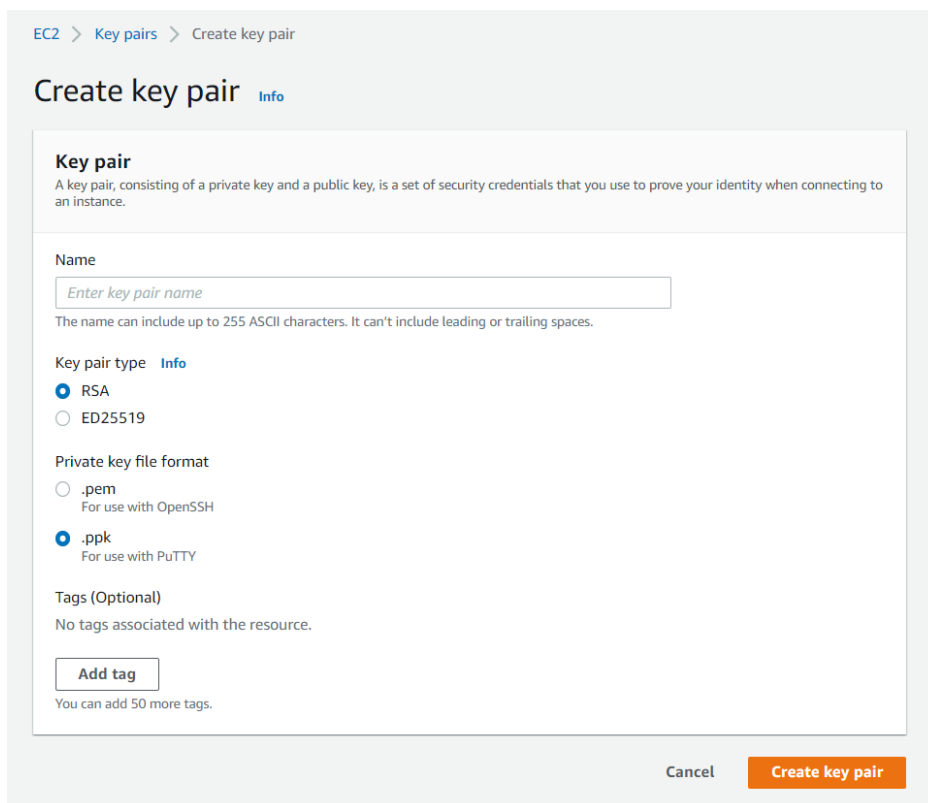
1. Problem Statement:

A tech companies is exploring ways to improve their sales. They want to start by analysing customer behaviour and gaining insights about product trends. Here, the role of big data analysts is among the most sought-after to gain the insights from abundance of data. Objective Through this assignment, as a big data analyst, we will extract data and gather insights from a real-life data set of an e-commerce company. We will analyse and gain insights about the clickstream data from a website so that we can extract insights about the customers behaviour.

2. Steps Involved:

2.1. Create Key Pair in AWS :

Open EC2 Management Console and select Create Key Pair. Enter a name for Key Pair, keeping the type as “RSA” and file format as “. ppk” and Select Create Key Pair. The Key Pair is downloaded.



The screenshot shows the 'Create key pair' page in the AWS Management Console. The breadcrumb trail at the top reads 'EC2 > Key pairs > Create key pair'. The main heading is 'Create key pair' with an 'Info' link. Below this is a section titled 'Key pair' with a description: 'A key pair, consisting of a private key and a public key, is a set of security credentials that you use to prove your identity when connecting to an instance.' The form contains the following fields and options:

- Name:** A text input field with the placeholder 'Enter key pair name'. Below it, a note states: 'The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.'
- Key pair type:** A dropdown menu with 'Info' link. The selected option is 'RSA' (indicated by a blue radio button). Another option is 'ED25519' (indicated by an empty radio button).
- Private key file format:** Two radio button options: '.pem' (For use with OpenSSH) and '.ppk' (For use with PuTTY). The '.ppk' option is selected.
- Tags (Optional):** A section with the text 'No tags associated with the resource.' and an 'Add tag' button. Below the button, it says 'You can add 50 more tags.'

At the bottom right of the form, there are two buttons: 'Cancel' and 'Create key pair' (which is orange and highlighted).

I have created a key pair with name Demo Keypair :

Key pairs (1/1) Info				
<input type="text" value="Filter key pairs"/>				
<input checked="" type="checkbox"/>	Name	Type	Fingerprint	ID
<input checked="" type="checkbox"/>	Demo_Keypair	rsa	7a:e5:cf:79:d6:e1:ce:8a:8f:2f:89:27:43:...	key-079981f8787149f69

2.2. Create S3 Bucket and upload the Data

Go to Amazon S3 > Select Create Bucket > Enter Bucket Name “upgradanuraganu” > Unselect “Block all public access” > Select Create Bucket.

Once the bucket is created then we will upload the files which is provided for the case study. We will click on our bucket and then click on upload to upload the files.

Amazon S3 > upgradanuraganu

upgradanuraganu [Info](#)

[Objects](#) [Properties](#) [Permissions](#) [Metrics](#) [Management](#) [Access Points](#)

Objects (2)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

[Refresh](#) [Copy S3 URI](#) [Copy URL](#) [Download](#) [Open](#) [Delete](#) [Actions](#) [Create folder](#) [Upload](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	2019-Nov.csv	csv	February 18, 2022, 08:44:42 (UTC+05:30)	520.6 MB	Standard
<input type="checkbox"/>	2019-Oct.csv	csv	February 18, 2022, 08:44:42 (UTC+05:30)	460.2 MB	Standard

Once the file is upload, we will select one object and copy the S3 URL. Same with the same object. Below is the object URL cleated in S3.

s3://upgradanuraganu/2019-Nov.csv

s3://upgradanuraganu/2019-Oct.csv

2.3. We will now create the EMR cluster :

Create EMR Cluster • Go to EMR > Select Create Cluster. Selected the General Configuration, Software Configuration, and Hardware Configuration as shown in below images. Click “Create Cluster

Software Configuration

Release

<input checked="" type="checkbox"/> Hadoop 2.10.1	<input type="checkbox"/> Zeppelin 0.10.0	<input type="checkbox"/> Livy 0.7.1
<input type="checkbox"/> JupyterHub 1.1.0	<input type="checkbox"/> Tez 0.9.2	<input type="checkbox"/> Flink 1.13.1
<input type="checkbox"/> Ganglia 3.7.2	<input type="checkbox"/> HBase 1.4.13	<input checked="" type="checkbox"/> Pig 0.17.0
<input checked="" type="checkbox"/> Hive 2.3.8	<input type="checkbox"/> Presto 0.261	<input type="checkbox"/> ZooKeeper 3.4.14
<input type="checkbox"/> JupyterEnterpriseGateway 2.1.0	<input type="checkbox"/> MXNet 1.8.0	<input type="checkbox"/> Sqoop 1.4.7
<input type="checkbox"/> Mahout 0.13.0	<input checked="" type="checkbox"/> Hue 4.9.0	<input type="checkbox"/> Phoenix 4.14.3
<input type="checkbox"/> Oozie 5.2.1	<input type="checkbox"/> Spark 2.4.8	<input type="checkbox"/> HCatalog 2.3.8
<input type="checkbox"/> TensorFlow 2.4.1		

Multiple master nodes (optional)

☐ Use multiple master nodes to improve cluster availability. [Learn more](#)

AWS Glue Data Catalog settings (optional)

☐ Use for Hive table metadata

Edit software settings

☒ Enter configuration ☐ Load JSON from S3

`classification=config-file-name,properties=[myKey1=myValue1,myKey2=myValue2]`

Steps (optional)

A step is a unit of work you submit to the cluster. For instance, a step might contain one or more Hadoop or Spark jobs. You can also submit additional steps to a cluster after it is running. [Learn more](#)

Concurrency: ☐ Run multiple steps at the same time to improve cluster utilization

After last step completes: ☒ Clusters enters waiting state

☐ Cluster auto-terminates

Step type

I have given the Cluster name as Demo_cluster

General Options

Cluster name

☒ Logging

S3 folder

☐ Log encryption

☒ Debugging

☐ Termination protection

Now, in next step we will choose the EC2 key pair which we have created in first step, my key pair is Demo Keypair. Then we will click on Create Cluster.

Security Options

EC2 key pair Demo_Keypair i

☒ Cluster visible to all IAM users in account i

Permissions i

☒ Default ☐ Custom

Use default IAM roles. If roles are not present, they will be automatically created for you with managed policies for automatic policy updates.

EMR role [EMR_DefaultRole](#) ☐ Use EMR_DefaultRole_V2 i

EC2 instance profile [EMR_EC2_DefaultRole](#) i

Auto Scaling role [EMR_AutoScaling_DefaultRole](#) i

Security Configuration

EC2 security groups

An EC2 security group acts as a virtual firewall for your cluster nodes to control inbound and outbound traffic. There are two types of security groups you can configure, [EMR managed security groups](#) and [additional security groups](#). EMR will [automatically update](#) the rules in the EMR managed security groups in order to launch a cluster. [Learn more](#).

Type	EMR managed security groups EMR will automatically update the selected group	Additional security groups EMR will not modify the selected groups
Master	Default: sg-0b895bffcaff5ac1b (ElasticMapReduce)	No security groups selected
Core & Task	Default: sg-0b8fe8810131c3f4b (ElasticMapReduce)	No security groups selected

[Create a security group](#)

Cancel

Previous

Create cluster

- Now, Cluster is in Starting phase which means then the cluster is going to start it is looking for the free system for provision.

Cluster: Demo_Cluster Starting

Summary Application user interfaces Monitoring Hardware Configurations Events Steps Bootstrap actions

Summary

ID: j-1AWPA8B0NG1YP
Creation date: 2022-02-22 16:32 (UTC+5:30)
Elapsed time: 0 seconds
After last step completes: Cluster waits
Termination protection: Off [Change](#)
Tags: -- [View All / Edit](#)
Master public DNS: --

Application user interfaces

Persistent user interfaces : --
On-cluster user interfaces : --

Security and access

Key name: Demo_Keypair
EC2 instance profile: EMR_EC2_DefaultRole
EMR role: EMR_DefaultRole
Auto Scaling role: EMR_AutoScaling_DefaultRole
Visible to all users: All [Change](#)
Security groups for Master: [sg-0b895bffcaff5ac1b](#) (ElasticMapReduce-master)
Security groups for Core & Task: [sg-0b8fe8810131c3f4b](#) (ElasticMapReduce-slave)

Configuration details

Release label: emr-5.34.0
Hadoop distribution: Amazon 2.10.1
Applications: Hive 2.3.8, Pig 0.17.0, Hue 4.9.0
Log URI: s3://aws-logs-452886565518-us-east-1/elasticmapreduce/
EMRFS consistent view: Disabled
Custom AMI ID: --

Network and hardware

Availability zone: --
Subnet ID: [subnet-036e8338d2f89e609](#)
Master: Provisioning 1 m4.large
Core: Provisioning 1 m4.large
Task: --
Cluster scaling: Not enabled
Auto-termination: Terminate if Idle for 1 hour

- Now the cluster is in waiting state which means it is running now we can perform the operation.

Clone
Terminate
AWS CLI export

Cluster: Demo_Cluster Waiting Cluster ready after last step completed.

Summary
Application user interfaces
Monitoring
Hardware
Configurations
Events
Steps
Bootstrap actions

Summary

ID: j-1AWPA8B0NG1YP


Creation date: 2022-02-22 16:32 (UTC+5:30)

Elapsed time: 14 minutes

After last step completes: Cluster waits

Termination protection: Off [Change](#)

Tags: -- [View All / Edit](#)

Master public DNS: ec2-3-237-25-234.compute-1.amazonaws.com 


[Connect to the Master Node Using SSH](#)

Configuration details

Release label: emr-5.34.0

Hadoop distribution: Amazon 2.10.1


Applications: Hive 2.3.8, Pig 0.17.0, Hue 4.9.0


Log URI: s3://aws-logs-452886565518-us-east-1/elasticmapreduce/ 

EMRFS consistent view: Disabled

Custom AMI ID: --


Application user interfaces

Persistent user interfaces : [YARN timeline server](#), [Tez UI](#)

On-cluster user interfaces : Not Enabled [Enable an SSH Connection](#)

Network and hardware

Availability zone: us-east-1f

Subnet ID: [subnet-036e8338d2f89e609](#) 

Master: Running 1 m4.large

Core: Running 1 m4.large

Task: --

Cluster scaling: Not enabled

Auto-termination: Terminate if idle for 1 hour

Security and access


Key name: Demo_Keypair


EC2 instance profile: EMR_EC2_DefaultRole

EMR role: EMR_DefaultRole

Auto Scaling role: EMR_AutoScaling_DefaultRole

Visible to all users: All [Change](#)

Security groups for Master: [sg-0b895bffc5ac1b](#)  (ElasticMapReduce-master)

Security groups for Core & Task: [sg-0b8fe8810131c3f4b](#)  (ElasticMapReduce-slave)

- We will open PuTTY to download PuTTY the steps to download is already provided, enable the SSH connection to open the CLI.

Step 1: Open an SSH Tunnel to the Amazon EMR Master Node - [Learn more](#)

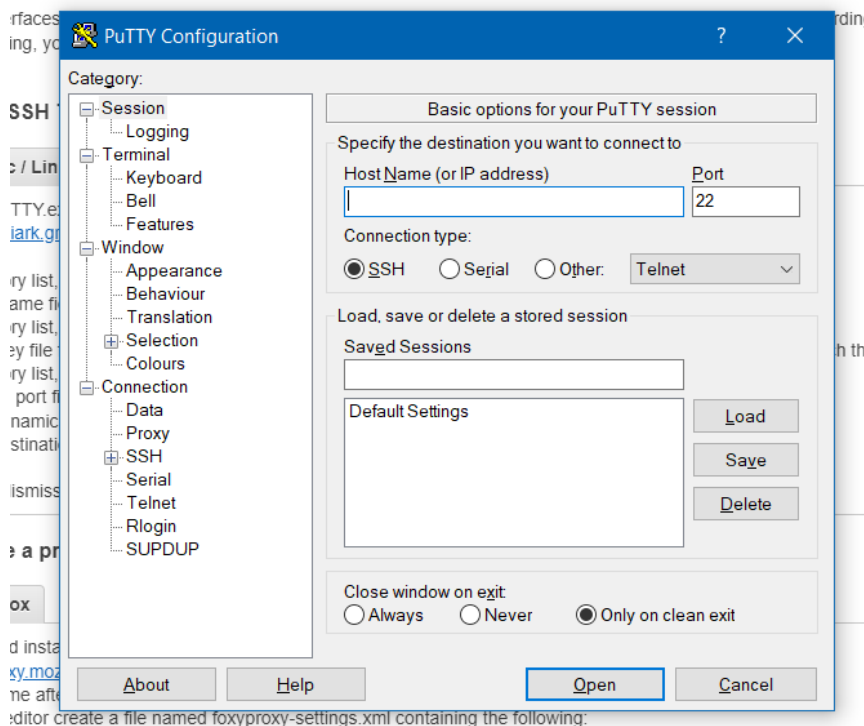
Windows Mac / Linux

1. Download PuTTY.exe to your computer from:
<http://www.chiark.greenend.org.uk/~sgtatham/putty/download.html>
2. Start PuTTY.
3. In the Category list, click Session
4. In the Host Name field, type **hadoop@ec2-3-237-25-234.compute-1.amazonaws.com**
5. In the Category list, expand Connection > SSH > Auth
6. For Private key file for authentication, click Browse and select the private key file (**Demo_Keypair.ppk**) used to launch the cluster.
7. In the Category list, expand Connection > SSH, and then click Tunnels.
8. In the Source port field, type **8157** (a randomly chosen, unused local port).
9. Select the Dynamic and Auto options.
10. Leave the Destination field empty and click Add.
11. Click Open.
12. Click Yes to dismiss the security alert.

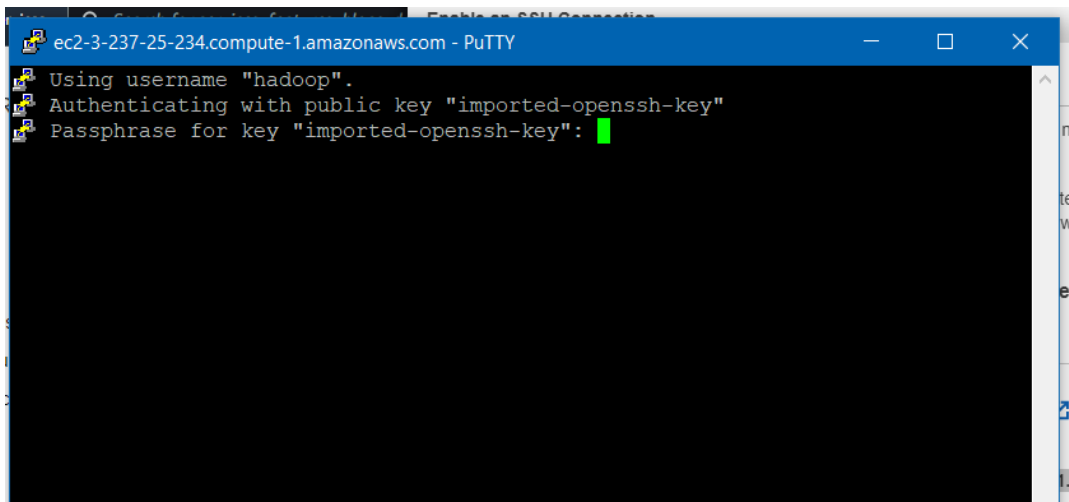
- Now, we will copy the HOST name **hadoop@ec2-3-237-25-234.compute-1.amazonaws.com** and paste it in PuTTY configuration in session tab.

Once cluster in running state we have to click on Master public DNS. Open the putty configuration and then give the host name (master node DNS) and then browse to the private key file location by clicking on Connection → SSH → Auth. Now we need to open PuTTY and connect to the master node by selecting the .ppk file.

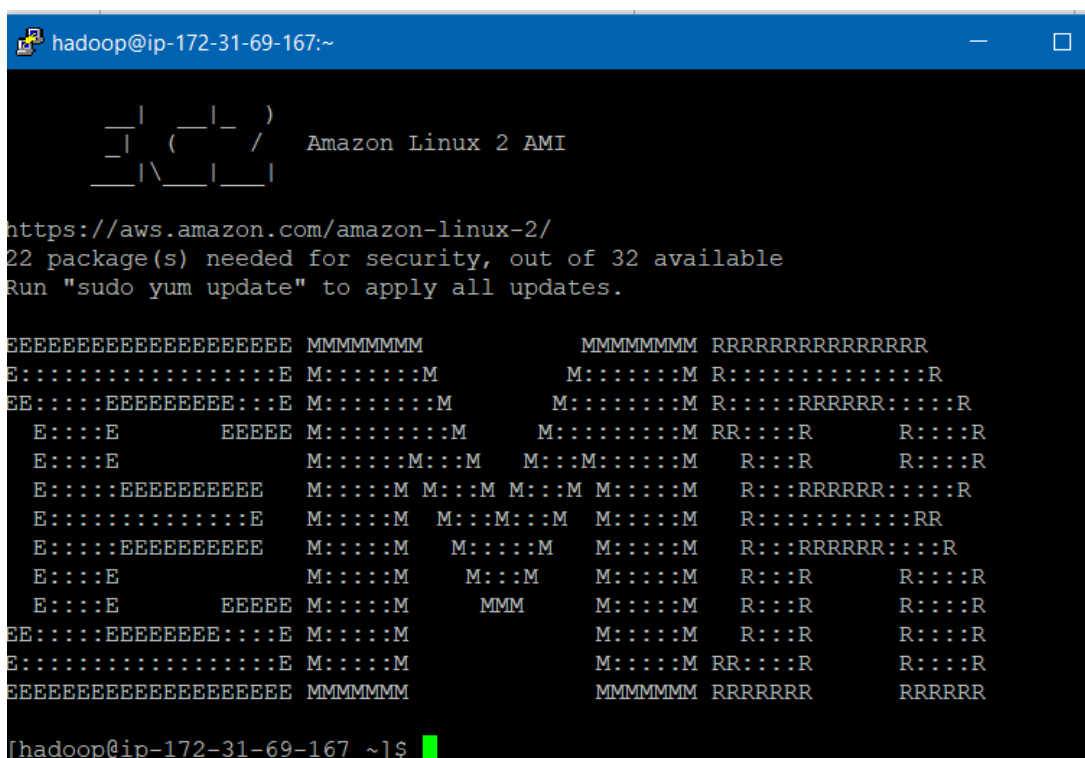
web server.



- Below is the security code which was enabled at the time of creation of private key pair.



EMR Opens, and below screen reflects that Connection to Hadoop is successful:



- Create a file directory :

Check the existing directories, and create a directory named 'Hive_Case_Study' in Hadoop. > 'Hive_Case_Study' is now created in Hadoop:

- Command for create a directory :

Query :

hadoop fs -mkdir /Hive_Case_Study

```
[hadoop@ip-172-31-69-167 ~]$ hadoop fs -mkdir / Hive_Case_Study
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/lib/hadoop/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLogger
Binder.class]
SLF4J: Found binding in [jar:file:/usr/share/aws/emr/emrfs/lib/slf4j-log4j12-1.7.12.jar!/org/slf4j/impl/Sta
ticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
```

- To view the directory :

Query:

hadoop fs -ls /

```
drwxr-xr-x - hadoop hdfsadmin group 0 2022-02-22 11:40 Hive_Case_Study
[hadoop@ip-172-31-69-167 ~]$
```

2.4. Move the data from the s3 buckets to the HDFS using the distributed copy command.

Loading the s3 public data set to created directory “Hive_Case_Study” in hadoop .

Query:

Hadoop distcp 's3://upgradanuraganu/2019-Oct.csv' / Hive_Case_Study/2019-Oct.csv

```
[hadoop@ip-172-31-69-167 ~]$ hadoop distcp 's3://upgradanuraganu/2019-Oct.csv' /Hive_Case_Study/2019-Oct.csv
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/lib/hadoop/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLogger
Binder.class]
SLF4J: Found binding in [jar:file:/usr/share/aws/emr/emrfs/lib/slf4j-log4j12-1.7.12.jar!/org/slf4j/impl/Sta
ticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
22/02/22 11:43:46 INFO tools.OptionsParser: parseChunkSize: blocksparchunk false
22/02/22 11:43:48 INFO tools.DistCp: Input Options: DistCpOptions{atomicCommit=false, syncFolder=false, del
eteMissing=false, ignoreFailures=false, overwrite=false, append=false, useDiff=false, useRdiff=false, fromS
napshot=null, toSnapshot=null, skipCRC=false, blocking=true, numListstatusThreads=0, maxMaps=20, mapBandwid
th=100, sslConfigurationFile='null', copyStrategy='uniformsize', preserveStatus=[], preserveRawXattrs=false
, atomicWorkPath=null, logPath=null, sourceFileListing=null, sourcePaths=[s3://upgradanuraganu/2019-Oct.csv
], targetPath=/Hive_Case_Study/2019-Oct.csv, targetPathExists=false, filtersFile='null', blocksPerChunk=0,
copyBufferSize=8192, verboseLog=false}
22/02/22 11:43:48 INFO client.RMPProxy: Connecting to ResourceManager at ip-172-31-69-167.ec2.internal/172.3
1.69.167:8032
22/02/22 11:43:48 INFO client.AHSPProxy: Connecting to Application History server at ip-172-31-69-167.ec2.in
ternal/172.31.69.167:10200
22/02/22 11:43:52 INFO tools.SimpleCopyListing: Paths (files+dirs) cnt = 1; dirCnt = 0
22/02/22 11:43:52 INFO tools.SimpleCopyListing: Build file listing completed.
```

```
DistCp Counters
  Bytes Copied=482542278
  Bytes Expected=482542278
  Files Copied=1
[hadoop@ip-172-31-69-167 ~]$
```

Hadoop distcp 's3://upgradanuraganu/2019-Nov.csv' /Hive_Case_Study/2019-Nov.csv

```
[hadoop@ip-172-31-69-167 ~]$ hadoop distcp 's3://upgradanuraganu/2019-Nov.csv' /Hive_Case_Study/2019-Nov.csv
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/lib/hadoop/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/share/aws/emr/emrfs/lib/slf4j-log4j12-1.7.12.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
22/02/22 11:49:53 INFO tools.OptionsParser: parseChunkSize: blocksperchunk false
22/02/22 11:49:54 INFO tools.DistCp: Input Options: DistCpOptions{atomicCommit=false, syncFolder=false, deleteMissing=false, ignoreFailures=false, overwrite=false, append=false, useDiff=false, useRdiff=false, fromSnapshot=null, toSnapshot=null, skipCRC=false, blocking=true, numListStatusThreads=0, maxMaps=20, mapBandwidth=100, sslConfigurationFile='null', copyStrategy='uniformsize', preserveStatus=[], preserveRawXattrs=false, atomicWorkPath=null, logPath=null, sourceFileListing=null, sourcePaths=[s3://upgradanuraganu/2019-Nov.csv], targetPath=Hive_Case_Study/2019-Nov.csv, targetPathExists=false, filtersFile='null', blocksPerChunk=0, copyBufferSize=8192, verboseLog=false}
22/02/22 11:49:54 INFO client.RMPProxy: Connecting to ResourceManager at ip-172-31-69-167.ec2.internal/172.31.69.167:8032
22/02/22 11:49:54 INFO client.AHSPProxy: Connecting to Application History server at ip-172-31-69-167.ec2.internal/172.31.69.167:10200
22/02/22 11:49:58 INFO tools.SimpleCopyListing: Paths (files+dirs) cnt = 1; dirCnt = 0
22/02/22 11:49:58 INFO tools.SimpleCopyListing: Build file listing completed.
22/02/22 11:49:58 INFO Configuration.deprecation: io.sort.mb is deprecated. Instead, use mapreduce.task.io.
```

```
2/02/22 11:50:00 INFO impl.YarnClientImpl: Submitted application application_1645528278761_0003
2/02/22 11:50:00 INFO mapreduce.Job: The url to track the job: http://ip-172-31-69-167.ec2.internal:20888/proxy/application_1645528278761_0003/
2/02/22 11:50:00 INFO tools.DistCp: DistCp job-id: job_1645528278761_0003
2/02/22 11:50:00 INFO mapreduce.Job: Running job: job_1645528278761_0003
2/02/22 11:50:09 INFO mapreduce.Job: Job job_1645528278761_0003 running in uber mode : false
2/02/22 11:50:09 INFO mapreduce.Job: map 0% reduce 0%
2/02/22 11:50:26 INFO mapreduce.Job: map 100% reduce 0%
2/02/22 11:50:30 INFO mapreduce.Job: Job job_1645528278761_0003 completed successfully
2/02/22 11:50:30 INFO mapreduce.Job: Counters: 38
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=224234
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=370
    HDFS: Number of bytes written=545839412
    HDFS: Number of read operations=11
```

```
Bytes Written=0
DistCp Counters
  Bytes Copied=545839412
  Bytes Expected=545839412
  Files Copied=1
[hadoop@ip-172-31-69-167 ~]$
```

- To check if both the files are copied from S3 to HDFS.

Query:

Hadoop fs -ls / Hive_case_study/

```
[hadoop@ip-172-31-69-167 ~]$ hadoop fs -ls / Hive_Case_Study/
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/lib/hadoop/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/share/aws/emr/emrfs/lib/slf4j-log4j12-1.7.12.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Found 5 items
drwxr-xr-x - hadoop hdfsadmingroup 0 2022-02-22 11:44 /Hive_Case_Study
drwxr-xr-x - hdfs hdfsadmingroup 0 2022-02-22 11:10 /apps
drwxrwxrwt - hdfs hdfsadmingroup 0 2022-02-22 11:12 /tmp
drwxr-xr-x - hdfs hdfsadmingroup 0 2022-02-22 11:10 /user
drwxr-xr-x - hdfs hdfsadmingroup 0 2022-02-22 11:10 /var
Found 2 items
-rw-r--r-- 1 hadoop hdfsadmingroup 545839412 2022-02-22 11:50 Hive_Case_Study/2019-Nov.csv
-rw-r--r-- 1 hadoop hdfsadmingroup 482542278 2022-02-22 11:57 Hive_Case_Study/2019-Oct.csv
[hadoop@ip-172-31-69-167 ~]$
```

2.5. Starting Hive:

Query :

```
[hadoop@ip-172-31-69-167 ~]$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/lib/hive/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/share/aws/emr/emrfs/lib/slf4j-log4j12-1.7.12.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive>
```

- We will create a database create a table and then load the data which was provided and execute the query to find/analyse the data which was provided.

1. CREATE DATABASE IF NOT EXISTS Hive_Case_Study;

```
hive> CREATE DATABASE IF NOT EXISTS Hive_Case_Study;
OK
Time taken: 0.046 seconds
```

2. Show databases; command to view the data base

```
hive> show databases;
OK
default
hive_case_study
Time taken: 0.087 seconds, Fetched: 2 row(s)
```

3. To describe the database command use is

Describe database Hive_Case_Study;

```
hive> describe database Hive_Case_Study;
OK
hive_case_study      hdfs://ip-172-31-69-167.ec2.internal:8020/user/hive/warehouse/hive_case_study.db h
adoop  USER
Time taken: 0.063 seconds, Fetched: 1 row(s)
hive>
```

4. To use the same database:

Query:

```
hive> use hive_case_study;
OK
Time taken: 0.056 seconds
hive>
```

- 2.6. Now the database is created we will now Create an external table "cosmeticDB":

Query:

```
CREATE EXTERNAL TABLE IF NOT EXISTS cosmeticDB (event_time timestamp, event_type
string, product_id string, category_id string, category_code string, brand string, price
decimal(10,3), user_id bigint, user_session string) ROW FORMAT SERDE
'org.apache.hadoop.hive.serde2.OpenCSVSerde' WITH SERDEPROPERTIES ("separatorChar" =
",", "quoteChar" = "\"", "escapeChar" = "\\") stored as textfile;
```

1. To view the table command used is (describe table name)

Query:

Describe cometicDB;

```

hive> CREATE EXTERNAL TABLE IF NOT EXISTS cosmeticDB (event_time timestamp, event_type string, product_id string, category_id string, category_code string, brand string, price decimal(10,3), user_id bigint, user_session string) ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' WITH SERDEPROPERTIES ("separatorChar" = ",", "quoteChar" = "\"", "escapeChar" = "\\") stored as textfile
> ;
OK
Time taken: 0.682 seconds
hive> describe cosmeticDB;
OK
event_time          string              from deserializer
event_type           string              from deserializer
product_id           string              from deserializer
category_id          string              from deserializer
category_code        string              from deserializer
brand                string              from deserializer
price                string              from deserializer
user_id              string              from deserializer
user_session         string              from deserializer
Time taken: 0.111 seconds, Fetched: 9 row(s)

```

Here if we observe all the data types are in 'string' by default. This is one of the limitations in serde. We need to cast the data types into desired ones.

2. Load the input Data into "cosmeticDB" table:

Query:

```
LOAD DATA INPATH 's3://upgradanuraganu/2019-Oct.csv' INTO TABLE cosmeticDB;
```

```
LOAD DATA INPATH 's3://upgradanuraganu/2019-Nov.csv' INTO TABLE cosmeticDB;
```

```

hive> load data inpath 's3://upgradanuraganu/2019-Oct.csv' into table cosmeticDB;
Loading data to table hive_case_study.cosmeticdb
OK
Time taken: 13.986 seconds
hive> load data inpath 's3://upgradanuraganu/2019-Nov.csv' into table cosmeticDB;
Loading data to table hive_case_study.cosmeticdb
OK
Time taken: 11.751 seconds
hive>

```

To check if the load data is correct for this, we will write a query to check the month(event_time) because the data is of 11(Nov) and 10(Oct).

```
hive> select * from cosmeticDB where month(event_time)=10 limit 5;
OK
2019-10-01 00:00:00 UTC cart 5773203 1487580005134238553 runail 2.62 46324
2019-10-01 00:00:03 UTC cart 5773353 1487580005134238553 runail 2.62 46324
2019-10-01 00:00:07 UTC cart 5881589 2151191071051219817 lovely 13.48 42968
2019-10-01 00:00:07 UTC cart 5723490 1487580005134238553 runail 2.62 46324
2019-10-01 00:00:15 UTC cart 5881449 1487580013522845895 lovely 0.56 42968
Time taken: 1.597 seconds, Fetched: 5 row(s)
hive> select * from cosmeticDB where month(event_time)=11 limit 5;
OK
2019-11-01 00:00:02 UTC view 5802432 1487580009286598681 0.32 56207
2019-11-01 00:00:09 UTC cart 5844397 1487580006317032337 2.38 55332
2019-11-01 00:00:10 UTC view 5837166 1783999064103190764 pnb 22.22 55613
2019-11-01 00:00:11 UTC cart 5876812 1487580010100293687 jessnail 3.16
2019-11-01 00:00:24 UTC remove_from_cart 5826182 1487580007483048900
Time taken: 0.265 seconds, Fetched: 5 row(s)
hive>
```

2.7. Applying Optimization Techniques Columns Partitioning and Bucketing:

1. Start with enabling the dynamic partitioning and bucketing.

```
hive> set hive.exec.dynamic.partition.mode=nonstrict;
hive> set hive.exec.dynamic.partition=true;
hive> set hive.enforce.bucketing=true;
hive>
```

2. Create a Table “cosmeticDB_P1”, PARTITIONED BY(event_type string) CLUSTERED BY (user_id) INTO 10 buckets o Create a Table “RetailDB_EC2”, PARTITIONED BY(event_type string) CLUSTERED BY (user_id) INTO 10 buckets

```
CREATE TABLE IF NOT EXISTS cosmeticDB_P1 (event_time timestamp, product_id string,
category_id string, category_code string, brand string, price float, user_id bigint,
user_session string) PARTITIONED BY (event_type string) CLUSTERED BY (user_id) INTO 10
BUCKETS ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS
TEXTFILE;
```

```
hive> CREATE TABLE IF NOT EXISTS cosmeticDB_P1 (event_time timestamp, product_id string, category_id string, category_code string, brand string, price float, user_id bigint, user_session string) PARTITIONED BY (event_type string) CLUSTERED BY (user_id) INTO 10 BUCKETS ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS TEXTFILE;
OK
Time taken: 0.155 seconds
```

3. Load the Data into the tables:

Insert into table cosmeticDB_P1 partition (event_type) select cast (replace (event_time, 'UTC', '') as timestamp), product_id, category_id, category_code, brand, cast(price as float), cast(user_id as bigint), user_session, event_type from cosmeticDB;

```
hive> Insert into table cosmeticDB_P1 partition (event_type) select cast (replace (event_time, 'UTC', '') as timestamp), product_id, category_id, category_code, brand, cast(price as float), cast(user_id as bigint), user_session, event_type from cosmeticDB;
Query ID = hadoop_20220224211735_6eb829bf-92ec-4915-8268-84d3cd287ac1
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1645734550072_0006)

Map 1: 0/8      Reducer 2: 0/10
Map 1: 0/8      Reducer 2: 0/10
Map 1: 0/8      Reducer 2: 0/10
Map 1: 0(+2)/8  Reducer 2: 0/10
Map 1: 0(+3)/8  Reducer 2: 0/10
Map 1: 0(+3)/8  Reducer 2: 0/10
Map 1: 0(+3)/8  Reducer 2: 0/10
Map 1: 0(+3)/8  Reducer 2: 0/10
Map 1: 0(+3)/8  Reducer 2: 0/10
Map 1: 0(+3)/8  Reducer 2: 0/10
Map 1: 0(+3)/8  Reducer 2: 0/10
Map 1: 0(+3)/8  Reducer 2: 0/10
Map 1: 0(+3)/8  Reducer 2: 0/10
Map 1: 0(+3)/8  Reducer 2: 0/10
Map 1: 0(+3)/8  Reducer 2: 0/10
```

```
Time taken to load dynamic partitions: 0.633 seconds
Time taken for adding to write entity : 0.003 seconds
OK
Time taken: 0.636 seconds
```

3. Question and Answer:

Q1. Find the total revenue generated due to purchases made in October.

Query:

```
SELECT sum(price) FROM cosmeticDB WHERE month(event_time)=10 and
event_type='purchase';
```

Normal Table : (Without Partitioning)

```
hive> select sum(price) from cosmeticDB where month(event_time)=10 and event_type='purchase';
Query ID = hadoop_20220222124619_e80e36db-93e8-4c7d-8384-7c92198ec322
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1645528278761_0007)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    8         8         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02  [=====>>>] 100%  ELAPSED TIME: 73.97 s
-----
OK
1211538.429999982
```

Partition Table :

```
SELECT sum(price) FROM cosmeticDB_P1 where month(event_time)=10 and
event_time='purchase';
```


If we compare the query optimization of partitioned table it is almost 43 second's less than normal table.

```
hive> SELECT sum(price) FROM cosmeticDB_P1 where month(event_time)=10 and event_type='purchase';
Query ID = hadoop_20220222135419_4a44a0a7-5ceb-4ae5-98b3-f7999dd35cd9
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1645528278761_0012)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container    SUCCEEDED   7         7           0         0         0         0
Reducer 2 ..... container  SUCCEEDED   1         1           0         0         0         0
-----
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 29.06 s
-----
OK
1211538.4300000072
Time taken: 37.148 seconds, Fetched: 1 row(s)
```

Q2. Write a query to yield the total sum of purchases per month in a single output.

Query:

```
SELECT MONTH(event_time) AS Month, COUNT(event_type) AS Purchases FROM
cosmeticDB_P1 WHERE event_type = 'purchase' GROUP BY MONTH(event_time);
```

```
hive> SELECT MONTH(event_time) AS Month, COUNT(event_type) AS Purchases FROM cosmeticDB_P1 WHERE
event_type = 'purchase' GROUP BY MONTH(event_time);
Query ID = hadoop_20220222141049_39589188-d40e-4a5a-ab69-eea8a88cdcee
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1645528278761_0013)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container    SUCCEEDED   7         7           0         0         0         0
Reducer 2 ..... container  SUCCEEDED   2         2           0         0         0         0
-----
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 29.84 s
-----
OK
10      245624
11      322417
Time taken: 30.473 seconds, Fetched: 2 row(s)
hive>
```

Q3. Find the total revenue generated due to purchases made in October.

Query:

```
SELECT sum (case when month(event_time)=10 then price else -1*price end) as
change_in_revenue FROM cosmeticDB_P1 WHERE month(event_time) in (10,11) and
```

event_type='purchase';

```
hive> SELECT sum (case when month(event_time)=10 then price else -1*price end) as change_in_r
Query ID = hadoop_20220224213251_6b1f2528-8bb3-451f-b940-eaf8356983d4
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1645734550072_0006)

-----
      VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    7        7        0        0        0        0
Reducer 2 ..... container  SUCCEEDED    1        1        0        0        0        0
-----
VERTICES: 02/02  [=====>>] 100%  ELAPSED TIME: 34.98 s
-----
OK
-319478.4699999841
Time taken: 35.946 seconds, Fetched: 1 row(s)
hive>
```

Q4. Find distinct categories of products. Categories with null category code can be ignored.

Query :

select distinct category_code as product_category from cosmeticDB_P1 where
category_code <> ''

```
hive> select distinct(category_code) as product_category from cosmeticDB_P1 where category_code <> '';
Query ID = hadoop_20220224215932_4a1a3928-f9ef-44ad-a634-d5f076bc7a81
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1645734550072_0006)

-----
      VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    6        6        0        0        0        0
Reducer 2 ..... container  SUCCEEDED    1        1        0        0        0        0
-----
VERTICES: 02/02  [=====>>] 100%  ELAPSED TIME: 63.18 s
-----
OK
accessories.bag
accessories.cosmetic_bag
apparel.glove
appliances.environment.air_conditioner
appliances.environment.vacuum
appliances.personal.hair_cutter
category_code
furniture.bathroom.bath
furniture.living_room.cabinet
furniture.living_room.chair
sport.diving
stationery.cartridge
Time taken: 63.785 seconds, Fetched: 12 row(s)
hive>
```

Q5. Find the total number of products available under each category.

Query :

Select category_code as category,count(product_id) as products from cosmeticDB_P1 where
category_code <> '' group by category_code;

```

hive> Select category_code as category,count(product_id) as products from cosmeticDB_P1 where category_code <> ''
group by category_code;
Query ID = hadoop_20220224215401_2fda4fec-96ed-4668-956e-4195a3edfb88
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1645734550072_0006)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    6         6         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 63.50 s
-----
OK
accessories.bag 11681
accessories.cosmetic_bag      1248
apparel.glove 18232
appliances.environment.air_conditioner 332
appliances.environment.vacuum 59761
appliances.personal.hair_cutter 1643
category_code 2
furniture.bathroom.bath 9857
furniture.living_room.cabinet 13439
furniture.living_room.chair 308
sport.diving 2
stationery.cartridge 26722
Time taken: 64.11 seconds, Fetched: 12 row(s)
hive>

```

Q6. Which brand had the maximum sales in October and November combined?

Query :

Select brand, sum(price) as max_sales from cosmeticDB_P1 where brand <> '' and event_type = 'purchase' group by brand order by max_sales desc limit 1;

```

hive> Select brand, sum(price) as max_sales from cosmeticDB_P1 where brand <> '' and event_type = 'purchase' group by brand order by max_sales desc limit 1;
Query ID = hadoop_20220224215103_945668d6-dbe3-4dfd-a66b-39c9e188764a
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1645734550072_0006)

Map 1: 0/7      Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 0/7      Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 0(+1)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 0(+2)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 0(+3)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 0(+3)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 0(+3)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 0(+3)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 0(+3)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 1(+3)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 2(+3)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 3(+3)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 4(+3)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 5(+2)/7  Reducer 2: 0/2      Reducer 3: 0/1
Map 1: 5(+2)/7  Reducer 2: 0(+1)/2      Reducer 3: 0/1
Map 1: 6(+1)/7  Reducer 2: 0(+2)/2      Reducer 3: 0/1
Map 1: 7/7      Reducer 2: 0(+2)/2      Reducer 3: 0/1
Map 1: 7/7      Reducer 2: 1(+1)/2      Reducer 3: 0(+1)/1
Map 1: 7/7      Reducer 2: 2/2      Reducer 3: 0(+1)/1
Map 1: 7/7      Reducer 2: 2/2      Reducer 3: 1/1
OK
runail 148297.93999999957
Time taken: 26.816 seconds, Fetched: 1 row(s)
hive>

```

Q7. Which brands increased their sales from October to November?

Query:

With diff_brand as (SELECT brand, month(event_time) as Month, sum(price) as sales , dense_rank() over(partition by brand order by sum(price) desc) as rank FROM cosmeticDB_P1 where brand <> '' and event_type= 'purchase' GROUP BY brand, month(event_time) ORDER BY brand, Month)

SELECT brand from diff_brand where rank =1 and Month= 11

```
hive> With diff_brand as(SELECT brand, month(event_time) as Month, sum(price) as sales , dense_rank() over(partition
by brand order by sum(price) desc) as rank FROM cosmeticDB_P1 where brand <> '' and event_type= 'purchase' GROUP BY b
rand, month(event_time) ORDER BY brand, Month
> )SELECT brand from diff_brand where rank =1 and Month= 11;
Query ID = hadoop_20220224222939_fe79bd4d-bec5-489c-aa52-2596e7e80314
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1645734550072_0007)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   7       7          0        0        0        0
Reducer 2 ..... container  SUCCEEDED   2       2          0        0        0        0
Reducer 3 ..... container  SUCCEEDED   2       2          0        0        0        0
Reducer 4 ..... container  SUCCEEDED   1       1          0        0        0        0
-----
VERTICES: 04/04  [=====>>>] 100%  ELAPSED TIME: 29.77 s
-----
```

```
airnails
art-visage
artex
aura
balbcare
barbie
batiste
beautix
beauty-free
beautyblender
beauugreen
benovy
binacil
bioaqua
biore
blixz
bluesky
bodyton
bpw.style
browxenna
candy
carmex
chi
coifin
concept
cosima
cosmoprofi
cristalinas
cutrin
de.lux
deoproce
depilflax
```

dewal
dizao
domix
ecocraft
ecolab
egomania
elizavecca
ellips
elskin
enjoy
entity
eos
estel
estelare
f.o.x
farmavita
farmona
fedua
finish
fly
foamie
freedecor
freshbubble
gehwoi
glysolid
godefroy

greymy
happyfons
haruyama
helloganic
igrobeauty
ingarden
inm
insight
irisk
italwax
jaguar
jas
jessnail
joico
juno
kaaral
kamill
kapous
kares
kaypro
keen
kerasys
kims
kinetics
kiss
kocostar
koelcia
koelf
konad
kosmekka
laboratorium
lador
ladykin
latinoil
levissime
levrana
lianail
likato
limoni
lovely
lowence
mane
marathon
markell
marutaka-foot
masura
matreshka
matrix
mavala
metzger

```
milv
miskin
missha
moyou
nagaraku
naomi
nefertiti
neoleor
nirvel
nitrile
oniq
orly
osmo
ovale
plazan
polarus
profepil
profhenna
protokeratin
provoc
rasyan
refectocil
rosi
roubloff
runail
s.care
sanoto
severina
shary
shik
skinity
skinlite
smart
soleo
solomeya
sophin
staleks
strong
supertan
swarovski
tertio
treaclemoon
trind
uno
uskusi
veraclara
vilenta
yoko
yu-r
zeitun
Time taken: 30.541 seconds, Fetched: 160 row(s)
hive>
```

Q8. Your company wants to reward the top 10 users of its website with a Golden Customer plan. Write a query to generate a list of top 10 users who spend the most.

Query:

```
select user_id,sum(price) as purchase, dense_rank() over( order by sum(price) desc) as rank
from cosmeticDB_P1 where event_type= 'purchase' group by user_id limit 10;
```

```
Time taken: 92.305 seconds, Fetched: 10 row(s)
hive> select user_id,sum(price) as purchase, dense_rank() over( order by sum(price)
> desc) as rank from cosmeticDB_P1 group by user_id limit 10;
Query ID = hadoop_20220224223644_a0b96bb1-e054-4df1-9ce0-9f3abec39c04
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1645734550072_0008)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	6	6	0	0	0	0	
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0	
Reducer 3	container	SUCCEEDED	6	6	0	0	0	0	

```
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 84.29 s
OK
557616099      63266.96999999997      1
557956487      52370.219999999965     2
550388516      46264.280000000264     3
531900924      43504.710000000014     4
352394658      28205.90999999997      5
550353491      25317.260000000013     6
443045778      23742.680000000008     7
479928991      23540.599999999999     8
554848397      23359.430000000008     9
526213023      22983.280000000024     10
Time taken: 92.305 seconds, Fetched: 10 row(s)
hive>
```

Observations:

1. For less query time we should always partition the table and also make bucket.
2. In our data set also we observed that performance rate increased when we used the partitioned table.
3. In event_type table compared to all the other event types purchase is less than other two event types.
4. Highest number of products available under appliances. environment.vaccume category.
5. The total revenue is high in November month than October month.
6. Runail brand has highest sales compared with other brands.
7. The user_id: 557616099 spent most compared to others and selected for reward and recognition program.

4. Dropping the Database and Terminating the Cluster

For dropping the database :

Query:

Show database;

Drop database hive_case_study cascade; (we use cascade to delete the table which is created in database. Without that it will give an error.)

```
hive> show databases;
OK
default
hive_case_study
Time taken: 0.022 seconds, Fetched: 2 row(s)
hive> drop database hive_case_study;
FAILED: Execution Error, return code 1 from org.apache.hadoop.hive.ql.exec.DDLTask. InvalidOperationException(message:Database hive_case_study is not empty)
hive> drop database hive_case_study cascade;
OK
Time taken: 0.703 seconds
hive> show databases;
OK
default
Time taken: 0.028 seconds, Fetched: 1 row(s)
hive>
```

Quitting Hive from PuTTY :

```
OK
default
Time taken: 0.028 seconds, Fetched: 1 row(s)
hive> quit;
[hadoop@ip-172-31-71-86 ~]$
```

Terminate EMR Cluster

- Select Terminate > Select Terminate in the pop-up window >Status changes to Terminating.

Clone

Terminate

AWS CLI export

Cluster: Demo_Cluster
Waiting
Cluster ready after last step completed.

Summary

Application user interfaces

Monitoring

Hardware

Configurations

Events

Steps

Bootstrap actions

Summary

ID: j-33PX3MQJS8QGZ
Creation date: 2022-02-25 01:49 (UTC+5:30)
Elapsed time: 2 hours, 33 minutes
After last step completes: Cluster waits
Termination protection: Off [Change](#)
Tags: -- [View All](#) / [Edit](#)
Master public DNS: ec2-34-204-197-146.compute-1.amazonaws.com [Connect to the Master Node Using SSH](#)

Configuration details

Release label: emr-5.34.0
Hadoop distribution: Amazon 2.10.1
Applications: Hive 2.3.8, Pig 0.17.0, Hue 4.9.0
Log URI: s3://aws-logs-452886565518-us-east-1/elasticmapreduce/ [View Log](#)
EMRFS consistent view: Disabled
Custom AMI ID: --

Application user interfaces

Persistent user interfaces [View](#): [YARN timeline server](#), [Tez UI](#)
On-cluster user interfaces [View](#): Not Enabled [Enable an SSH Connection](#)

Network and hardware

Availability zone: us-east-1f
Subnet ID: [subnet-036e8338d2f89e609](#) [View Subnet](#)
Master: Running 1 m4.large
Core: Running 1 m4.large
Task: --
Cluster scaling: Not enabled
Auto-termination: Terminate if idle for 1 hour

Cluster: Demo_Cluster
Terminating
Terminated by user request

Summary

Application user interfaces

Monitoring

Hardware

Configurations

Events

Steps

Bootstrap actions

Summary

ID: j-33PX3MQJS8QGZ
Creation date: 2022-02-25 01:49 (UTC+5:30)
Elapsed time: 2 hours, 33 minutes
After last step completes: Cluster waits
Termination protection: Off
Tags: --
Master public DNS: ec2-34-204-197-146.compute-1.amazonaws.com [Connect to the Master Node Using SSH](#)

Configuration details

Release label: emr-5.34.0
Hadoop distribution: Amazon 2.10.1
Applications: Hive 2.3.8, Pig 0.17.0, Hue 4.9.0
Log URI: s3://aws-logs-452886565518-us-east-1/elasticmapreduce/ [View Log](#)
EMRFS consistent view: Disabled
Custom AMI ID: --

Application user interfaces

Persistent user interfaces [View](#): [YARN timeline server](#), [Tez UI](#)
On-cluster user interfaces [View](#): Not Enabled [Enable an SSH Connection](#)

Network and hardware

Availability zone: us-east-1f
Subnet ID: [subnet-036e8338d2f89e609](#) [View Subnet](#)
Master: Running 1 m4.large
Core: Running 1 m4.large
Task: --
Cluster scaling: Not enabled
Auto-termination: Terminate if idle for 1 hour

Cluster is successfully terminated :

Clone Terminate AWS CLI export

Cluster: Demo_Cluster Terminated Terminated by user request

Summary Application user interfaces Monitoring Hardware Configurations Events Steps Bootstrap actions

Summary

ID: j-33PX3MQJS8QGZ
Creation date: 2022-02-25 01:49 (UTC+5:30)
End date: 2022-02-25 04:26 (UTC+5:30)
Elapsed time: 2 hours, 36 minutes
After last step completes: Cluster waits
Termination protection: Off
Tags: --
Master public DNS: ec2-34-204-187-146.compute-1.amazonaws.com

Configuration details

Release label: emr-5.34.0
Hadoop distribution: Amazon 2.10.1
Applications: Hive 2.3.8, Pig 0.17.0, Hue 4.9.0
Log URI: s3://aws-logs-452886565518-us-east-1/elasticmapreduce/
EMRFS consistent view: Disabled
Custom AMI ID: --

Made By :

Anurag Aditya & Hazarathaiah G